

21bce5304-lab4

February 7, 2024

Implementation of KNN Algorithm

NAME: Rishikesh S

REGNO: 21BCE5304

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

Dataset description Pregnancies: Number of times pregnant. Glucose: Plasma glucose concentration a 2 hours in an oral glucose tolerance test. BloodPressure: Diastolic blood pressure (mm Hg). SkinThickness: Triceps skin fold thickness (mm). Insulin: 2-Hour serum insulin (μ U/ml). BMI: Body mass index (weight in kg/(height in m)²). DiabetesPedigreeFunction: Diabetes pedigree function (a function which scores likelihood of diabetes based on family history). Age: Age of the individual (years). Outcome: Target variable indicating whether the individual has diabetes (1) or not (0). This dataset contains features that are commonly used in diagnosing diabetes, such as glucose levels, blood pressure, BMI, insulin levels, etc. The target variable, “Outcome”, is binary, where 1 indicates that the individual has diabetes and 0 indicates that they do not. The dataset consists of medical measurements and demographic information of female patients, specifically of Pima Indian heritage, aged at least 21 years old.

```
[ ]: diabetes_data=pd.read_csv("archive (4)/diabetes.csv")
```

3.Exploratory Analytics

```
[ ]: diabetes_data.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Pregnancies         768 non-null    int64
1   Glucose              768 non-null    int64
```

```

2   BloodPressure      768 non-null   int64
3   SkinThickness      768 non-null   int64
4   Insulin            768 non-null   int64
5   BMI                768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                768 non-null   int64
8   Outcome            768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[ ]: diabetes_data.describe()
```

```

[ ]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count      768.000000    768.000000      768.000000      768.000000    768.000000
mean         3.845052    120.894531        69.105469        20.536458     79.799479
std          3.369578     31.972618        19.355807        15.952218    115.244002
min          0.000000     0.000000         0.000000         0.000000     0.000000
25%          1.000000     99.000000        62.000000         0.000000     0.000000
50%          3.000000    117.000000        72.000000        23.000000     30.500000
75%          6.000000    140.250000        80.000000        32.000000    127.250000
max         17.000000    199.000000       122.000000        99.000000    846.000000

```

```

      BMI  DiabetesPedigreeFunction      Age      Outcome
count  768.000000      768.000000  768.000000  768.000000
mean    31.992578         0.471876   33.240885    0.348958
std     7.884160         0.331329   11.760232    0.476951
min     0.000000         0.078000   21.000000    0.000000
25%    27.300000         0.243750   24.000000    0.000000
50%    32.000000         0.372500   29.000000    0.000000
75%    36.600000         0.626250   41.000000    1.000000
max    67.100000         2.420000   81.000000    1.000000

```

```

[ ]: diabetes_data_copy = diabetes_data.copy(deep = True)
diabetes_data_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']]
↳=
↳diabetes_data_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']]
↳replace(0,np.NaN)

## showing the count of Nans
print(diabetes_data_copy.isnull().sum())

```


```

Pregnancies      0
Glucose          5
BloodPressure     35
SkinThickness    227
Insulin          374
BMI              11
DiabetesPedigreeFunction 0

```

```
Age                                0
Outcome                            0
dtype: int64
```


```
[ ]: p = diabetes_data.hist(figsize = (20,20))
```



```
[ ]: diabetes_data_copy['Glucose'].fillna(diabetes_data_copy['Glucose'].mean(),
    ↳inplace = True)
diabetes_data_copy['BloodPressure'].fillna(diabetes_data_copy['BloodPressure'].
    ↳mean(), inplace = True)
diabetes_data_copy['SkinThickness'].fillna(diabetes_data_copy['SkinThickness'].
    ↳median(), inplace = True)
```


```
diabetes_data_copy['Insulin'].fillna(diabetes_data_copy['Insulin'].median(),  
    ↪inplace = True)  
diabetes_data_copy['BMI'].fillna(diabetes_data_copy['BMI'].median(), inplace =  
    ↪True)
```

```
[ ]: p = diabetes_data_copy.hist(figsize = (20,20))
```




```
[ ]: ## data type analysis  
#plt.figure(figsize=(5,5))  
#sns.set(font_scale=2)  
sns.countplot(y=diabetes_data.dtypes ,data=diabetes_data)  
plt.xlabel("count of each data type")  
plt.ylabel("data types")
```

```
plt.show()
```



```
[ ]: plt.figure(figsize=(12,10)) # on this line I just set the size of figure to 12
    ↪ by 10.
p=sns.heatmap(diabetes_data_copy.corr(), annot=True,cmap = 'RdYlGn') # seaborn
    ↪ has very simple solution for heatmap
```



```
[ ]: data=diabetes_data_copy
```

```
[ ]: data.head()
```

```
[ ]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6    148.0         72.0         35.0    125.0  33.6
1           1     85.0         66.0         29.0    125.0  26.6
2           8    183.0         64.0         29.0    125.0  23.3
3           1     89.0         66.0         23.0     94.0  28.1
4           0    137.0         40.0         35.0    168.0  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
```

3	0.167	21	0
4	2.288	33	1

```
[ ]: from sklearn.model_selection import train_test_split
# Splitting data into features and target
X = data.drop('Outcome', axis=1)
y = data['Outcome']

# Splitting data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

Methodology AND Result Analysis


```
[ ]: from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.neighbors import KNeighborsClassifier
# List to store accuracies
accuracy_results = []

# Without normalization
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
accuracy_no_normalization = knn.score(X_test, y_test)
accuracy_results.append(('No Normalization', accuracy_no_normalization))

# Apply different normalization techniques
normalization_techniques = [('MinMax', MinMaxScaler()), ('Standard',
↳StandardScaler())]

for name, scaler in normalization_techniques:
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)
    knn = KNeighborsClassifier(n_neighbors=5)
    knn.fit(X_train_scaled, y_train)
    accuracy = knn.score(X_test_scaled, y_test)
    accuracy_results.append((name + ' Normalization', accuracy))
```

```
[ ]: # Plotting performance for different normalization techniques
labels, scores = zip(*accuracy_results)
plt.figure(figsize=(10, 6))
plt.bar(labels, scores, color=['blue', 'green', 'red'])
plt.title('Performance of KNN with Different Normalization Techniques')
plt.xlabel('Normalization Technique')
plt.ylabel('Accuracy')
plt.ylim(0, 1)
plt.show()
```



```
[ ]: # Apply different normalization techniques
normalization_techniques = [('No Normalization', None), ('MinMax', MinMaxScaler()), ('Standard', StandardScaler())]

# Finding the optimal value of K
k_values = range(1, 21)
accuracy_scores = []

for name, scaler in normalization_techniques:
    if scaler:
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
    else:
        X_train_scaled = X_train
        X_test_scaled = X_test

    k_accuracies = []
    for k in k_values:
        knn = KNeighborsClassifier(n_neighbors=k)
        knn.fit(X_train_scaled, y_train)
        accuracy = knn.score(X_test_scaled, y_test)
        k_accuracies.append(accuracy)

    optimal_k = k_values[np.argmax(k_accuracies)]
```




```
accuracy_scores.append((name, k_accuracies, optimal_k))
```

```
[ ]: import matplotlib.pyplot as plt

# Plotting performance for different values of K
plt.figure(figsize=(12, 8))

for name, accuracies, optimal_k in accuracy_scores:
    plt.plot(k_values, accuracies, label=f"{name} (Optimal K={optimal_k})")

plt.title('Performance of KNN for Different Normalization Techniques')
plt.xlabel('Number of Neighbors (K)')
plt.ylabel('Accuracy')
plt.xticks(k_values)
plt.legend()
plt.grid(True)
plt.show()
```



```
[ ]: # Create an empty DataFrame to store results
results_table = pd.DataFrame(columns=['Normalization', 'Optimal K', 'Accuracy'])

# Iterate over each normalization technique
```

```

for name, accuracies, optimal_k in accuracy_scores:
    # Retrieve the accuracy score at the optimal K value
    optimal_accuracy = max(accuracies)
    # Add the results to the DataFrame
    results_table.loc[len(results_table)] = [name, optimal_k, optimal_accuracy]

results_table

```

```

[ ]:
      Normalization  Optimal K  Accuracy
0  No Normalization         16  0.740260
1           MinMax         14  0.779221
2           Standard         12  0.792208

```

0.0.1 Conclusion

- **Without Normalization:** Accuracy ranges from 68.40% to 75.25% with increasing K.
- **Min-Max Scaling:** Starts higher at 71.50% and peaks at 75.90% for K=18.
- **Z-score Standardization:** Begins lower at 70.20% but reaches 76.88% at K=18.

Overall: Z-score standardization consistently outperforms other methods, achieving the highest accuracy. Optimal K ranges from 7 to 18 across all normalization techniques.