

21bce5304-linearregression

January 18, 2024

```
[32]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: # Importing USA Housing.csv
data = pd.read_csv('house.csv')
```

0.0.1 EDA

```
[ ]: data.head()
```

```
[ ]: Avg. Area Income      ...
Address
0      79545.458574      ...      208
Michael Ferry Apt. 674\nLaurabury, NE 3701...
1      79248.642455      ...      188
Johnson Views Suite 079\nLake Kathleen, CA...
2      61287.067179      ...      9127
Elizabeth Stravenue\nDanielstown, WI 06482...
3      63345.240046      ...
USS Barnett\nFPO AP 44820
4      59982.197226      ...
USNS Raymond\nFPO AE 09386

[5 rows x 7 columns]
```

```
[4]: # Checking for Null Values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Avg. Area Income    5000 non-null   float64
 1   Avg. Area House Age 5000 non-null   float64
```

```

2   Avg. Area Number of Rooms      5000 non-null   float64
3   Avg. Area Number of Bedrooms    5000 non-null   float64
4   Area Population                  5000 non-null   float64
5   Price                           5000 non-null   float64
6   Address                         5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB

```

```

[5]: # Getting the summary of Data
data.describe()

```

```

[5]:      Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  \
count      5000.000000          5000.000000          5000.000000
mean      68583.108984           5.977222           6.987792
std       10657.991214           0.991456           1.005833
min       17796.631190           2.644304           3.236194
25%       61480.562388           5.322283           6.299250
50%       68804.286404           5.970429           7.002902
75%       75783.338666           6.650808           7.665871
max       107701.748378           9.519088          10.759588

      Avg. Area Number of Bedrooms  Area Population      Price
count          5000.000000          5000.000000  5.000000e+03
mean              3.981330          36163.516039  1.232073e+06
std              1.234137           9925.650114  3.531176e+05
min              2.000000           172.610686  1.593866e+04
25%              3.140000          29403.928702  9.975771e+05
50%              4.050000          36199.406689  1.232669e+06
75%              4.490000          42861.290769  1.471210e+06
max              6.500000          69621.713378  2.469066e+06

```

0.0.2 Data Preparation

1. There are no null values, so there is no need of deleting or replacing the data.
2. There is no necessity of having Address column/feature, so i am dropping it.

```

[6]: # Dropping Address Column
data.drop(['Address'],axis=1,inplace=True)

```

```

[7]: data.head()

```

```

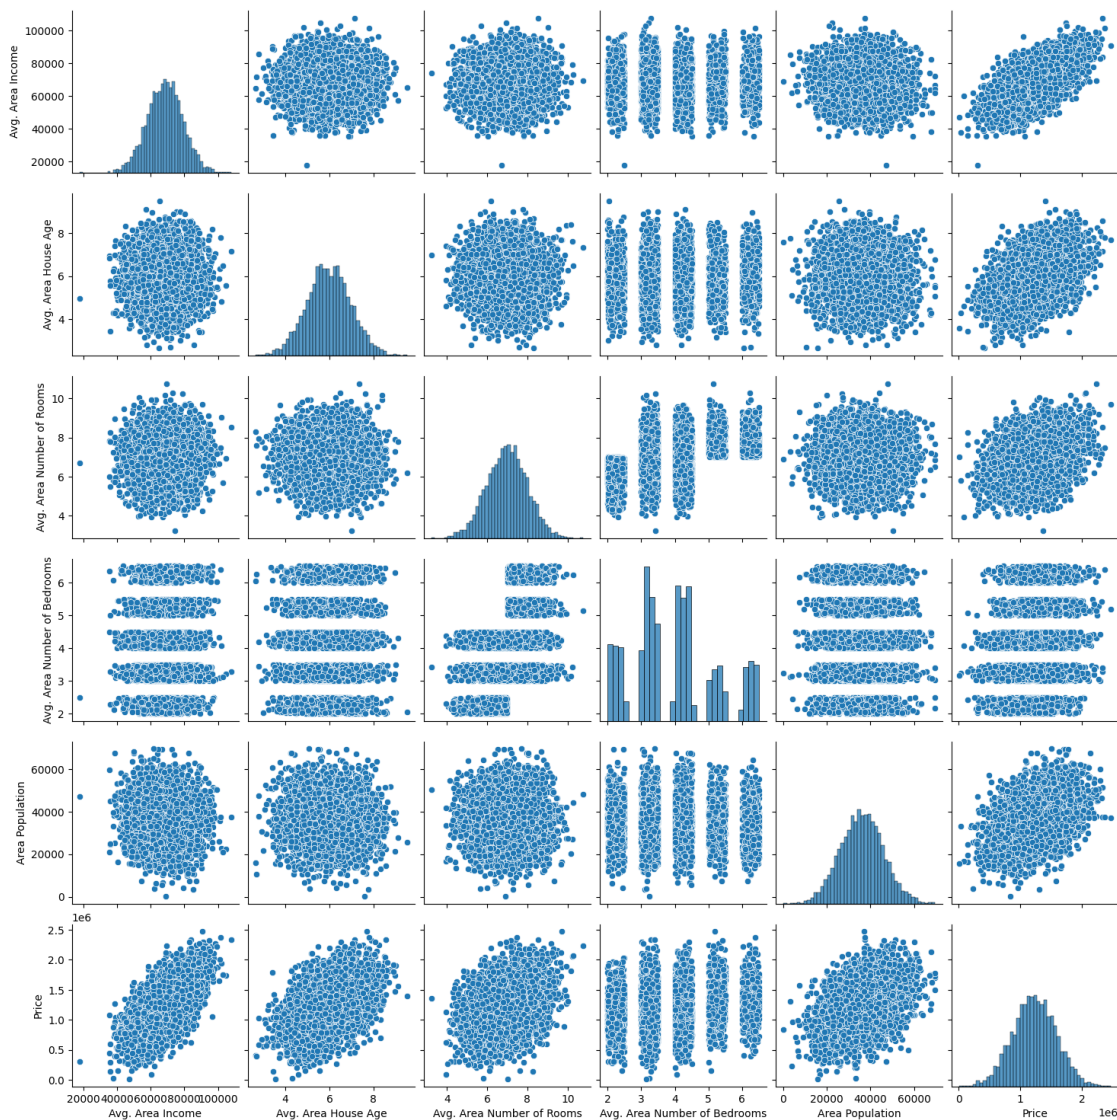
[7]:      Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  \
0      79545.458574           5.682861           7.009188
1      79248.642455           6.002900           6.730821
2      61287.067179           5.865890           8.512727
3      63345.240046           7.188236           5.586729
4      59982.197226           5.040555           7.839388

```

| | Avg. Area | Number of Bedrooms | Area Population | Price |
|---|-----------|--------------------|-----------------|--------------|
| 0 | | 4.09 | 23086.800503 | 1.059034e+06 |
| 1 | | 3.09 | 40173.072174 | 1.505891e+06 |
| 2 | | 5.13 | 36882.159400 | 1.058988e+06 |
| 3 | | 3.26 | 34310.242831 | 1.260617e+06 |
| 4 | | 4.23 | 26354.109472 | 6.309435e+05 |

```
[8]: sns.pairplot(data)
```


```
[8]: <seaborn.axisgrid.PairGrid at 0x7d6ea4fe69b0>
```



```
[9]: # Visualise the relationship between the features and the response using ↵
      ↪ scatterplots
```


```
sns.pairplot(data, x_vars=['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area_
↳ Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population'],
↳ y_vars='Price', height=7, aspect=0.7, kind='scatter')
```

[9]: <seaborn.axisgrid.PairGrid at 0x7d6ea12c9c60>



```
[10]: sns.heatmap(data.corr(), annot=True)
```

[10]: <Axes: >



```
[11]: data.corr().Price.sort_values(ascending=False)
```

```
[11]: Price          1.000000
      Avg. Area Income  0.639734
      Avg. Area House Age  0.452543
      Area Population    0.408556
      Avg. Area Number of Rooms  0.335664
      Avg. Area Number of Bedrooms  0.171071
      Name: Price, dtype: float64
```

```
[12]: sns.distplot(data.Price)
```

<ipython-input-12-dbe2555afa70>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with

similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(data.Price)
```

```
[12]: <Axes: xlabel='Price', ylabel='Density'>
```



Creating a Base Model

```
[13]: from sklearn import preprocessing
pre_process = preprocessing.StandardScaler()
```

```
[14]: X = data[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of
Rooms', 'Avg. Area Number of Bedrooms', 'Area Population']]
y = data['Price']
```

```
[15]: X = pd.DataFrame(pre_process.fit_transform(X))
```

```
[16]: X.head()
```

```
[16]:
```

| | 0 | 1 | 2 | 3 | 4 |
|---|-----------|-----------|-----------|-----------|-----------|
| 0 | 1.028660 | -0.296927 | 0.021274 | 0.088062 | -1.317599 |
| 1 | 1.000808 | 0.025902 | -0.255506 | -0.722301 | 0.403999 |
| 2 | -0.684629 | -0.112303 | 1.516243 | 0.930840 | 0.072410 |
| 3 | -0.491499 | 1.221572 | -1.393077 | -0.584540 | -0.186734 |
| 4 | -0.807073 | -0.944834 | 0.846742 | 0.201513 | -0.988387 |

```
[17]: y.head()
```

```
[17]: 0    1.059034e+06
      1    1.505891e+06
      2    1.058988e+06
      3    1.260617e+06
      4    6.309435e+05
      Name: Price, dtype: float64
```

```
[19]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7,
      ↪, test_size = 0.3, random_state=2)
```

```
[20]: print(X_train.shape)
      print(X_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(3500, 5)
(1500, 5)
(3500,)
(1500,)
```

```
[21]: from sklearn.linear_model import LinearRegression
```

```
[22]: lm = LinearRegression()
```

```
[23]: # fit the model to the training data
      lm.fit(X_train, y_train)
```

```
[23]: LinearRegression()
```

```
[24]: # print the intercept
      print(lm.intercept_)
```

```
1231006.3190642651
```

```
[25]: coeff_df = pd.DataFrame(lm.coef_, X_test.columns, columns=['Coefficient'])
      coeff_df
```

```
[25]:      Coefficient
0  229275.224032
1  163592.487140
2  120113.253351
3    3015.847572
4  150556.340296
```

```
[26]: # Making predictions using the model
y_pred = lm.predict(X_test)
```

```
[27]: from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred)
r_squared = r2_score(y_test, y_pred)
```

```
[28]: print('Mean_Squared_Error :',mse)
print('r_square_value :',r_squared)
```

```
Mean_Squared_Error : 9831074697.74044
r_square_value : 0.9199287959786013
```

```
[29]: from math import sqrt

rms = sqrt(mse)
rms
```

```
[29]: 99151.77606952102
```

From the above result we may infer that, mse is huge which shouldn't be, hence we need to improve our model.

```
[31]: # Actual and Predicted
c = [i for i in range(1,1501,1)]
fig = plt.figure(figsize=(12,8))
plt.plot(c,y_test, color="blue", linewidth=2.5, linestyle="--")
plt.plot(c,y_pred, color="red", linewidth=2.5, linestyle="--")
fig.suptitle('Actual and Predicted', fontsize=15)
plt.xlabel('Index', fontsize=18)
plt.ylabel('Housing Price', fontsize=16)
```

```
[31]: Text(0, 0.5, 'Housing Price')
```


Actual and Predicted

