

21bce5304-multipleregression

January 21, 2024

Objective: Implementation of Multiple Linear Regression

NAME: Rishikesh S

REGNO: 21BCE5304

```
[61]: import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import numpy as np
```

2.Dataset Description

The dataset you're referring to is related to temperature estimation.

```
[81]: data=pd.read_csv("startup.csv")
```

3.Exploratory Analytics

```
[82]:
```

```
data.describe()
```

```
[82]:
```

	R&D Spend	Administration	Marketing Spend	Profit
count	50.000000	50.000000	50.000000	50.000000
mean	73721.615600	121344.639600	211025.097800	
			112012.639200	
std	45902.256482	28017.802755	122290.310726	
				40306.180338
min	0.000000	51283.140000	0.000000	14681.400000
25%	39936.370000	103730.875000	129300.132500	
				90138.902500
50%	73051.080000	122699.795000	212716.240000	
			107978.190000	
75%	101602.800000	144842.180000	299469.085000	
			139765.977500	
max	165349.200000	182645.560000	471784.100000	
			192261.830000	

```
[83]: data.head()
```

```
[83]: R&D Spend Administration Marketing Spend State Profit 0
165349.20 136897.80 471784.10 New York 192261.83
1 162597.70 151377.59 443898.53 California 191792.06
```

```

2  153441.51      101145.55  407934.54  Florida 191050.39
3  144372.41      118671.85  383199.62  New York 182901.99
4  142107.34       91391.77  366168.42  Florida 166187.94

```

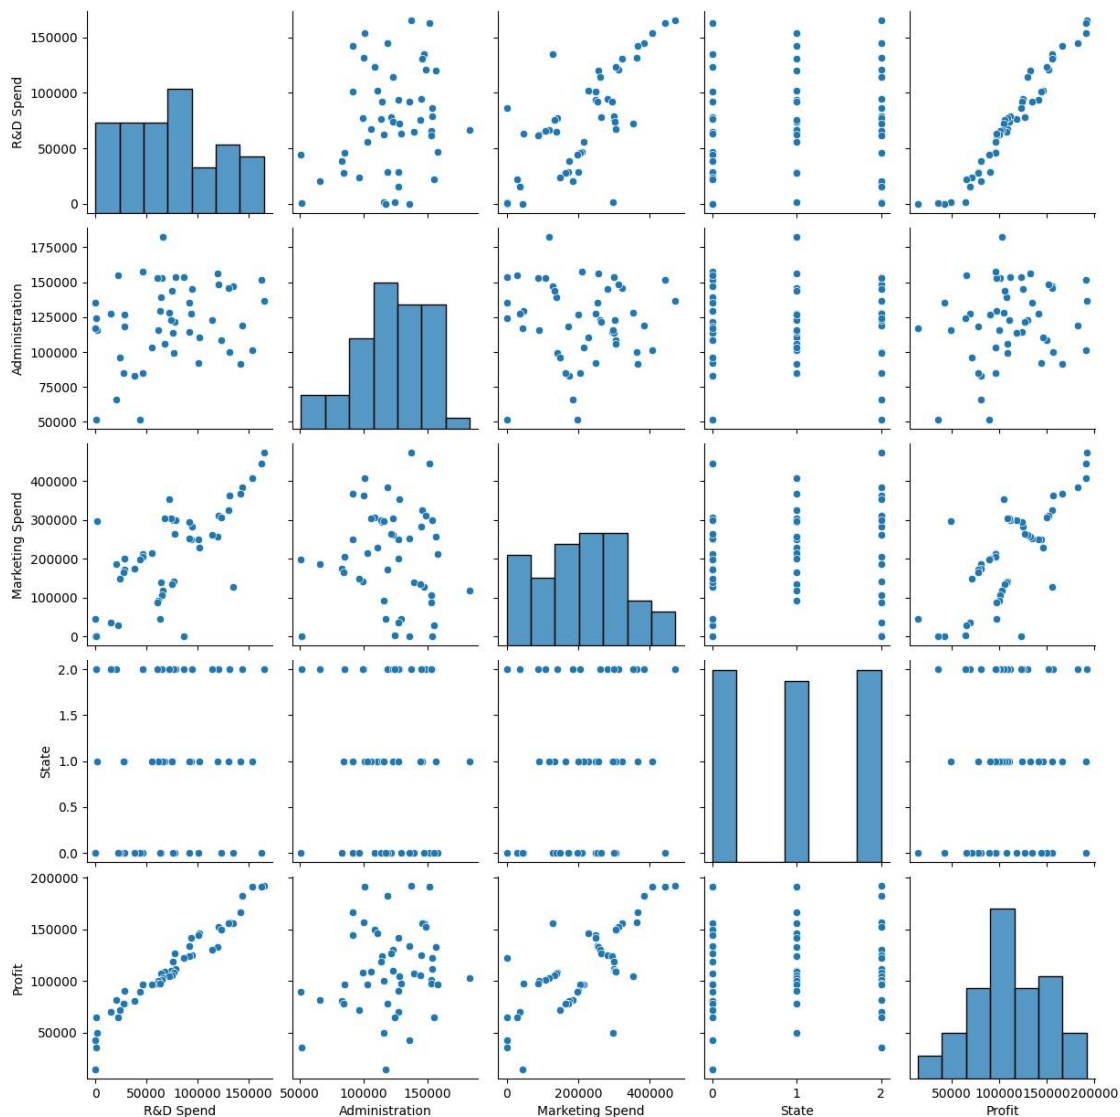
```
[84]: data["State"].unique()
```

```
[84]: array(['New York', 'California', 'Florida'], dtype=object)
```

```
[85]: from sklearn.preprocessing import LabelEncoder
      lbl=LabelEncoder()
      data["State"]=lbl.fit_transform(data["State"])
```

```
[86]: sns.pairplot(data)
```

```
[86]: <seaborn.axisgrid.PairGrid at 0x7fab66add20>
```



```
[87]: sns.distplot(data['Profit'])
```

```
<ipython-input-87-5c9dc59bcd4>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

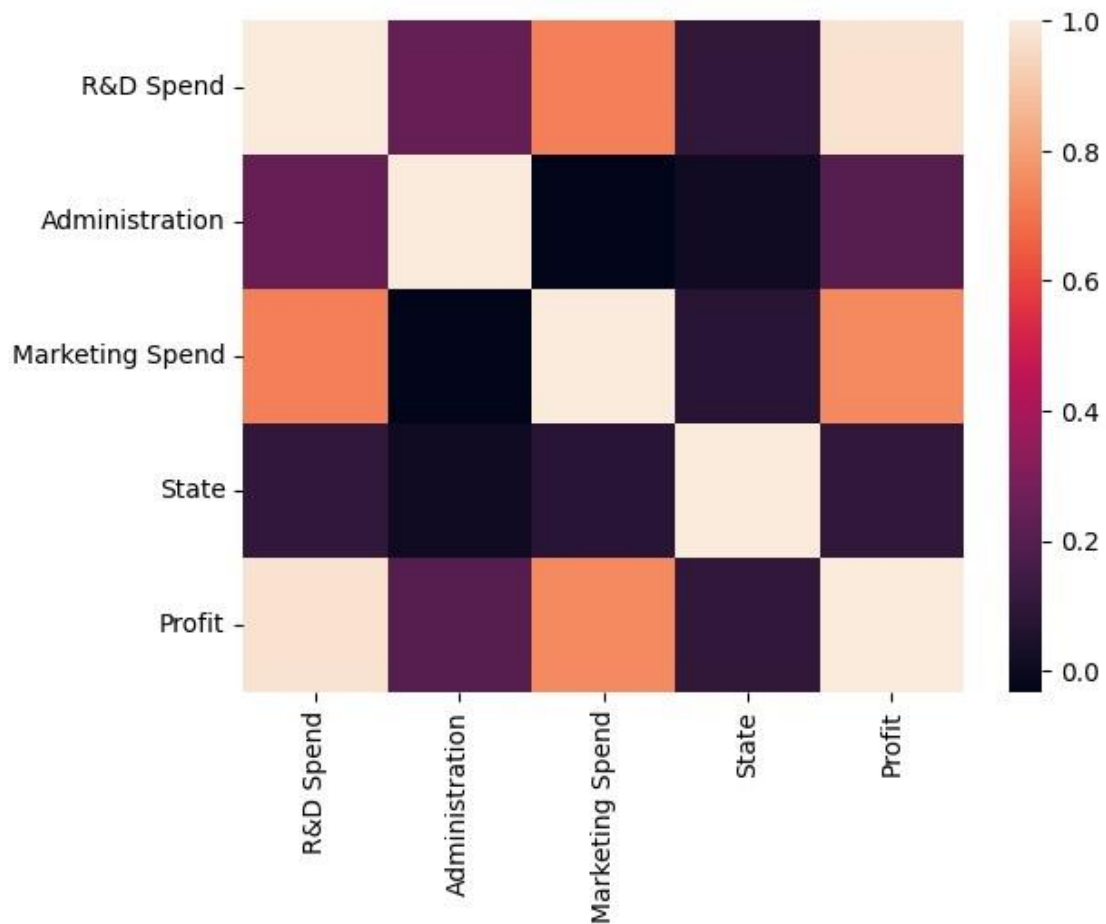
```
sns.distplot(data['Profit'])
```

```
[87]: <Axes: xlabel='Profit', ylabel='Density'>
```



```
[88]: sns.heatmap(data.corr())
```

```
[88]: <Axes: >
```



```
[89]: x = data.iloc[:, :-1].values
      y=data.iloc[:, -1].values
```

```
[90]: print(x)
      print(y)
```

```
[ [1.6534920e+05 1.3689780e+05 4.7178410e+05 2.0000000e+00]
  [1.6259770e+05 1.5137759e+05 4.4389853e+05 0.0000000e+00]
  [1.5344151e+05 1.0114555e+05 4.0793454e+05 1.0000000e+00]
  [1.4437241e+05 1.1867185e+05 3.8319962e+05 2.0000000e+00]
  [1.4210734e+05 9.1391770e+04 3.6616842e+05 1.0000000e+00]
  [1.3187690e+05 9.9814710e+04 3.6286136e+05 2.0000000e+00]
  [1.3461546e+05 1.4719887e+05 1.2771682e+05 0.0000000e+00]
  [1.3029813e+05 1.4553006e+05 3.2387668e+05 1.0000000e+00]
  [1.2054252e+05 1.4871895e+05 3.1161329e+05 2.0000000e+00]
  [1.2333488e+05 1.0867917e+05 3.0498162e+05 0.0000000e+00]
  [1.0191308e+05 1.1059411e+05 2.2916095e+05 1.0000000e+00]
```

```

[1.0067196e+05 9.1790610e+04 2.4974455e+05 0.0000000e+00]
[9.3863750e+04 1.2732038e+05 2.4983944e+05 1.0000000e+00]
[9.1992390e+04 1.3549507e+05 2.5266493e+05 0.0000000e+00]
[1.1994324e+05 1.5654742e+05 2.5651292e+05 1.0000000e+00]
[1.1452361e+05 1.2261684e+05 2.6177623e+05 2.0000000e+00]
[7.8013110e+04 1.2159755e+05 2.6434606e+05 0.0000000e+00]
[9.4657160e+04 1.4507758e+05 2.8257431e+05 2.0000000e+00]
[9.1749160e+04 1.1417579e+05 2.9491957e+05 1.0000000e+00]
[8.6419700e+04 1.5351411e+05 0.0000000e+00 2.0000000e+00]
[7.6253860e+04 1.1386730e+05 2.9866447e+05 0.0000000e+00]
[7.8389470e+04 1.5377343e+05 2.9973729e+05 2.0000000e+00]
[7.3994560e+04 1.2278275e+05 3.0331926e+05 1.0000000e+00]
[6.7532530e+04 1.0575103e+05 3.0476873e+05 1.0000000e+00]
[7.7044010e+04 9.9281340e+04 1.4057481e+05 2.0000000e+00]
[6.4664710e+04 1.3955316e+05 1.3796262e+05 0.0000000e+00]
[7.5328870e+04 1.4413598e+05 1.3405007e+05 1.0000000e+00]
[7.2107600e+04 1.2786455e+05 3.5318381e+05 2.0000000e+00]
[6.6051520e+04 1.8264556e+05 1.1814820e+05 1.0000000e+00]
[6.5605480e+04 1.5303206e+05 1.0713838e+05 2.0000000e+00]
[6.1994480e+04 1.1564128e+05 9.1131240e+04 1.0000000e+00]
[6.1136380e+04 1.5270192e+05 8.8218230e+04 2.0000000e+00]
[6.3408860e+04 1.2921961e+05 4.6085250e+04 0.0000000e+00]
[5.5493950e+04 1.0305749e+05 2.1463481e+05 1.0000000e+00]
[4.6426070e+04 1.5769392e+05 2.1079767e+05 0.0000000e+00]
[4.6014020e+04 8.5047440e+04 2.0551764e+05 2.0000000e+00]
[2.8663760e+04 1.2705621e+05 2.0112682e+05 1.0000000e+00]
[4.4069950e+04 5.1283140e+04 1.9702942e+05 0.0000000e+00]
[2.0229590e+04 6.5947930e+04 1.8526510e+05 2.0000000e+00]
[3.8558510e+04 8.2982090e+04 1.7499930e+05 0.0000000e+00]
[2.8754330e+04 1.1854605e+05 1.7279567e+05 0.0000000e+00]
[2.7892920e+04 8.4710770e+04 1.6447071e+05 1.0000000e+00]
[2.3640930e+04 9.6189630e+04 1.4800111e+05 0.0000000e+00]
[1.5505730e+04 1.2738230e+05 3.5534170e+04 2.0000000e+00]
[2.2177740e+04 1.5480614e+05 2.8334720e+04 0.0000000e+00]
[1.0002300e+03 1.2415304e+05 1.9039300e+03 2.0000000e+00]
[1.3154600e+03 1.1581621e+05 2.9711446e+05 1.0000000e+00]
[0.0000000e+00 1.3542692e+05 0.0000000e+00 0.0000000e+00]
[5.4205000e+02 5.1743150e+04 0.0000000e+00 2.0000000e+00]
[0.0000000e+00 1.1698380e+05 4.5173060e+04 0.0000000e+00]]
[192261.83 191792.06 191050.39 182901.99 166187.94 156991.12
156122.51
155752.6 152211.77 149759.96 146121.95 144259.4 141585.52 134307.35
132602.65 129917.04 126992.93 125370.37 124266.9 122776.86 118474.03
111313.02 110352.25 108733.99 108552.04 107404.34 105733.54
105008.31 103282.38 101004.64 99937.59 97483.56 97427.84 96778.92
96712.8

```

```
96479.51 90708.19 89949.14 81229.06 81005.76 78239.91 77798.83
71498.49 69758.98 65200.33 64926.08 49490.75 42559.73 35673.41
14681.4 ]
```

```
[43]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
random_state = 1)
```

```
[44]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test=sc.fit_transform(x_test)
```

```
[91]: data.head()
```

```
[91]:   R&D Spend  Administration  Marketing Spend  State      Profit
0  165349.20      136897.80   471784.10    2  192261.83
1  162597.70      151377.59   443898.53    0  191792.06
2  153441.51      101145.55   407934.54    1  191050.39
3  144372.41      118671.85   383199.62    2  182901.99
4  142107.34       91391.77   366168.42    1  166187.94
```

```
[92]: data.rename(columns={'Marketing Spend': 'Marketing_Spend'}, inplace=True)
data.head()
```

```
[92]: R&D Spend  Administration  Marketing_Spend  State      Profit
0  165349.20      136897.80   471784.10    2  192261.83
1  162597.70      151377.59   443898.53    0  191792.06
2  153441.51      101145.55   407934.54    1  191050.39
3  144372.41      118671.85   383199.62    2  182901.99
4  142107.34       91391.77   366168.42    1  166187.94
```

```
[93]: data.rename(columns={'R&D Spend': 'RD_Spend'}, inplace=True)
data.head()
```

```
[93]: RD_Spend  Administration  Marketing_Spend  State      Profit
0  165349.20      136897.80   471784.10    2  192261.83
1  162597.70      151377.59   443898.53    0  191792.06
2  153441.51      101145.55   407934.54    1  191050.39
3  144372.41      118671.85   383199.62    2  182901.99
4  142107.34       91391.77   366168.42    1  166187.94
```

Methodology AND Multiple Model Analysis

```
[94]: import statsmodels.formula.api as smf  
      from sklearn.metrics import mean_squared_error, mean_absolute_error  
      import numpy as np
```

```
[95]: # Initialize variables to store the results  
      results = []
```

```

# Create a function to build a model and print summary
def build_model(features, data, model_number):
    formula = f'Profit ~ {" + ".join(features)}'
    model = smf.ols(formula=formula, data=data).fit()

    # Get model performance metrics
    y_pred = model.predict(data[features])
    mse = mean_squared_error(data['Profit'], y_pred)
    mae = mean_absolute_error(data['Profit'], y_pred)
    rmse = np.sqrt(mse)
    r_squared = model.rsquared

    # Append results to the list
    p_values = model.pvalues[1:] # Exclude intercept
    results.append([model_number, formula, features, mse, mae, rmse, r_squared,
↪p_values])

    # Print model summary
    print(f"Model {model_number} - {formula}")
    print("MSE:", mse)
    print("MAE:", mae)
    print("RMSE:", rmse)
    print("R-squared:", r_squared)
    print("P-values:")
    print(p_values)
    print(model.summary())
    print("\n")

```

```

[97]: # Consider all features initially
all_features = [ 'RD_Spend', 'Administration', 'Marketing_Spend', 'State' ]
# Set a threshold for p-value
p_value_threshold = 0.05

# Build models and print summary based on p-values
model_number = 1
while all_features:
    build_model(all_features, data, model_number)
    model_number += 1

    # Get p-values for all features in the current model
    p_values = smf.ols(formula=f'Profit ~ {" + ".join(all_features)}',
↪data=data).fit().pvalues[1:] # Exclude intercept

    # Identify the feature with the highest p-value
    max_p_value_feature = p_values.idxmax()
    max_p_value = p_values[max_p_value_feature]

```



```

# Check if the highest p-value is above the threshold
if max_p_value > p_value_threshold:
    # Remove the feature with the highest p-value
    all_features.remove(max_p_value_feature)
else:
    # Break the loop if all features have p-values below the threshold
    break

```

Model 1 - Profit ~ RD_Spend + Administration + Marketing_Spend + State

MSE: 78416791.01666646

MAE: 6468.105695552672

RMSE: 8855.325573724913 R-squared: 0.9507462044842656

P-values:

RD_Spend 8.249206e-22

Administration 6.056771e-01

Marketing_Spend 1.086131e-01

State dtype: 9.889988e-01
float64

OLS Regression Results

```

=====
=====
Dep. Variable:          Profit  R-squared:          0.951
Model:                  OLS  Adj. R-squared:        0.946
Method:                 Least Squares  F-statistic:    217.2
Date:                  Sun, 21 Jan 2024  Prob (F-statistic): 8.51e-29
Time:                  13:33:05  Log-Likelihood:      -525.39
No. Observations:      50  AIC:          1061.
Df Residuals:          45  BIC:          1070.
Df Model:              4

```

Covariance Type: nonrobust

```

=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	5.014e+04	6804.555	7.369	0.000	3.64e+04	

6.38e+04

RD_Spend	0.8058	0.046	17.609	0.000	0.714
0.898					
Administration	-0.0268	0.052	-0.520	0.606	-0.131
0.077					
Marketing_Spend	0.0272	0.017	1.637	0.109	-0.006
0.061					
State	-22.3206	1609.829	-0.014	0.989	-
3220.041				3264.682	

=====

=====

Omnibus:	14.864	Durbin-Watson:	1.282
Prob(Omnibus):	0.001	Jarque-Bera (JB):	21.542
Skew:	-0.949	Prob(JB):	2.10e-05
Kurtosis:	5.596	Cond. No.	1.44e+06

=====

=====

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.44e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Model 2 - Profit ~ RD_Spend + Administration + Marketing_Spend

MSE: 78417126.01913084

MAE: 6471.45039610481

RMSE: 8855.344489015142 R-squared: 0.9507459940683246

P-values:

RD_Spend 2.634968e-22

Administration 6.017551e-01

Marketing_Spend 1.047168e-01

dtype: float64

OLS Regression Results

=====

=====

Dep. Variable:	Profit	R-squared:	0.951
Model:	OLS	Adj. R-squared:	0.948
Method:	Least Squares	F-statistic:	296.0
Date:	Sun, 21 Jan 2024	Prob (F-statistic):	4.53e-30
Time:	13:33:05	Log-Likelihood:	-525.39
No. Observations:	50	AIC:	1059.

Df Residuals: 46 BIC: 1066.
Df Model: 3

Covariance Type: nonrobust

```
=====
=====
              coef    std err          t      P>|t|      [0.025
0.975] -----
-----
---
Intercept      5.012e+04  6572.353      7.626    0.000    3.69e+04
6.34e+04
RD_Spend        0.8057      0.045     17.846    0.000      0.715
0.897
Administration -0.0268      0.051     -0.526    0.602     -0.130
0.076
Marketing_Spend  0.0272      0.016      1.655    0.105     -0.006
0.060
=====
=====
Omnibus:            14.838  Durbin-Watson:           1.282
Prob(Omnibus):      0.001  Jarque-Bera (JB):          21.442
Skew:               -0.949  Prob(JB):           2.21e-05
Kurtosis:           5.586  Cond. No.           1.40e+06
=====
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.4e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Model 3 - Profit ~ RD_Spend + Marketing_Spend

MSE: 78887897.00648756

MAE: 6499.319940113646

RMSE: 8881.885892449169 R-

squared: 0.9504503015559763

P-values:

RD_Spend 6.040433e-24

Marketing_Spend 6.003040e-02

dtype: float64

OLS Regression Results

```
=====
=====
```

```

Dep. Variable:          Profit  R-squared:          0.950
Model:                  OLS  Adj. R-squared:        0.948
Method:                 Least Squares  F-statistic:    450.8
Date:                  Sun, 21 Jan 2024  Prob (F-statistic): 2.16e-
                                                                31
Time:                  13:33:05  Log-Likelihood:      -
                                                                525.54
No. Observations:      50  AIC:                    1057.
Df Residuals:          47  BIC:                    1063.
Df Model:              2

```

Covariance Type: nonrobust

```

=====
=====
coef      std err      t      P>|t|      [0.025
0.975] -----
-----
---
Intercept      4.698e+04  2689.933    17.464    0.000    4.16e+04
5.24e+04
RD_Spend       0.7966    0.041    19.266    0.000    0.713
0.880
Marketing_Spend 0.0299    0.016    1.927    0.060   -0.001
0.061
=====
=====
Omnibus:      14.677  Durbin-Watson:      1.257
Prob(Omnibus): 0.001  Jarque-Bera (JB):      21.161
Skew:         -0.939  Prob(JB):      2.54e-05
Kurtosis:     5.575  Cond. No.      5.32e+05
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.32e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Model 4 - Profit ~ RD_Spend

MSE: 85120931.32706906

MAE: 6910.984354579612

RMSE: 9226.100548285232 R-squared: 0.9465353160804393

P-values:

```
RD_Spend    3.500322e-32
dtype: float64
```

OLS Regression Results

```
=====
Dep. Variable:          Profit  R-squared:                0.947
Model:                  OLS  Adj. R-squared:              0.945
Method:                 Least Squares  F-statistic:        849.8
Date:                  Sun, 21 Jan 2024  Prob (F-statistic):  3.50e-
                                                                32
Time:                  13:33:06  Log-Likelihood:          -
                                                                527.44
No. Observations:      50  AIC:                          1059.
Df Residuals:          48  BIC:                          1063.
Df Model:               1
Covariance Type:       nonrobust
```

```
=====
               ===== coef   std err   t      P>|t| [0.025    0.975]
-----
Intercept  4.903e+04  2537.897   19.320   0.000   4.39e+04   5.41e+04
RD_Spend    0.8543    0.029   29.151   0.000    0.795    0.913
=====
```

```
=====
Omnibus:            13.727  Durbin-Watson:                1.116
Prob(Omnibus):      0.001  Jarque-Bera (JB):            18.536
Skew:               -0.911  Prob(JB):                  9.44e-05
Kurtosis:           5.361  Cond. No.:                  1.65e+05
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.65e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Multiple regression on the best model obtained above using Scikit Learn

```
[98]: # Create a DataFrame from the results list results_df =
pd.DataFrame(results, columns=['No', 'Regression Model name',
    'Independent variables chosen', 'MSE', 'MAE', 'RMSE', 'R-
square', 'P-values'])
```

```
[100]: # Choose the best model based on the lowest RMSE
best_model =
results_df.loc[results_df['RMSE'].idxmin()]

# Extract features and target variable for the
best model best_features = best_model['Independent
variables chosen']
X =
data[best_features]
y = data['Profit']

[101]: from sklearn.model_selection import
train_test_split from sklearn.linear_model
import LinearRegression
from sklearn.metrics import mean_squared_error,
mean_absolute_error import numpy as np

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Train a Linear Regression model using scikit-
learn
model_sklearn = LinearRegression()
model_sklearn.fit(X_train, y_train)

[101]: LinearRegression()
```

Results

```
[102]: # Make predictions on the test
set y_pred =
model_sklearn.predict(X_test)

# Evaluate the model performance on the test
set mse_test = mean_squared_error(y_test,
y_pred) mae_test =
mean_absolute_error(y_test, y_pred) rmse_test
= np.sqrt(mse_test)
r_squared_test = model_sklearn.score(X_test, y_test)

# Print model performance on the test set
print("Best Model -", best_model['Regression Model name'])
print("Test Set Metrics:")
```

```

print("MSE:", mse_test)
print("MAE:", mae_test)
print("RMSE:", rmse_test)
print("R-squared:", r_squared_test)

```

Best Model - Profit ~ RD_Spend + Administration + Marketing_Spend
+ State Test Set Metrics:
MSE: 59510962.80787997
MAE: 6077.363300620398
RMSE: 7714.334890830185
R-squared: 0.9265108109341951

```
[103]: results_df
```

```
[103]:
```

```
[103] results_df
```

	No	Regression Model name	Independent variables chosen	MSE	MAE	RMSE	R-square	P-values
0	1	Profit ~ RD_Spend + Administration + Marketing...	[RD_Spend]	7.841679e+07	6468.105696	8855.325574	0.950746	RD_Spend 8.249206e-22 Administration...
1	2	Profit ~ RD_Spend + Administration + Marketing...	[RD_Spend]	7.841713e+07	6471.450396	8855.344489	0.950746	RD_Spend 2.634968e-22 Administration...
2	3	Profit ~ RD_Spend + Marketing_Spend	[RD_Spend]	7.888790e+07	6499.319940	8881.885892	0.950450	RD_Spend 6.040433e-24 Marketing_Spen...
3	4	Profit ~ RD_Spend	[RD_Spend]	8.512093e+07	6910.984355	9226.100548	0.946535	RD_Spend 3.500322e-32 dtype: float64